


| Oct 08, 2025 16:33 | train_cifar10.json | Page 1/4 |
|---|--------------------|----------|
| <pre> /* I deleted all the training log (train loss) since there were too many and ju st leave the validation loss log It should be easier to see the train/validation loss graph to see the traini ng summary Test set result is at the very bottom To be honest, I think I can get better result with better blocks_per_stage a nd layer_depths (more like resnet-18 or resnet-34 that I used for CIFAR-100) but I didn't have time to re-train it. Started: 2025-10-04T23:02:32.478968Z Ended: 2025-10-05T03:43:35.804946Z */ {"settings": {"debug": false, "random_seed": 31415, "model": {"input_depth": 3, "layer_depths": [16, 32, 64], "blocks_per_stage": [3, 3, 3], "num_classes": 10}, "data": {"batch_size": 64, "validation_size": 5000, "shuffle_buffer": 10000}, " training": {"num_epochs": 40, "learning_rate": 0.001, "weight_decay": 0.0001, "l og_interval": 100, "final_test": false, "checkpoint_dir": "PosixPath('hw04/check points')"}, "plotting": {"output_dir": "PosixPath('hw04/artifacts')", "figsize": [5, 3], "dpi": 200}}, "event": "Settings loaded", "logger": "hw04", "level": "i nfo", "timestamp": "2025-10-04T23:01:52.118846Z"} {"event": "Variant folder /root/tensorflow_datasets/cifar10/3.0.2 has no dataset _info.json"} {"event": "Generating dataset cifar10 (/root/tensorflow_datasets/cifar10/3.0.2)" } {"event": "Downloading https://www.cs.toronto.edu/~kriz/cifar-10-binary.tar.gz i nto /root/tensorflow_datasets/downloads/cifar10/cs.toronto.edu_kriz_cifar-10-bin aryODHPTiJLh3oLcXirEISTO7dkzyKjRCuol6lV8Wc6C7s.tar.gz.tmp.elcbf29ba6684e91b56e41 462da23de0..."} {"event": "Done writing /root/tensorflow_datasets/cifar10/incomplete.4A5IR3_3.0. 2/cifar10-train.tfrecord*. Number of examples: 50000 (shards: [50000])"} {"event": "Done writing /root/tensorflow_datasets/cifar10/incomplete.4A5IR3_3.0. 2/cifar10-test.tfrecord*. Number of examples: 10000 (shards: [10000])"} {"event": "Creating a tf.data.Dataset reading 1 files located in folders: /root/ tensorflow_datasets/cifar10/3.0.2."} {"event": "Constructing tf.data.Dataset cifar10 for split train, from /root/tens orflow_datasets/cifar10/3.0.2"} {"event": "Load dataset info from /root/tensorflow_datasets/cifar10/3.0.2"} {"event": "Creating a tf.data.Dataset reading 1 files located in folders: /root/ tensorflow_datasets/cifar10/3.0.2."} {"event": "Constructing tf.data.Dataset cifar10 for split test, from /root/tens orflow_datasets/cifar10/3.0.2"} {"event": "Unable to initialize backend 'tpu': INTERNAL: Failed to open libtpu.s o: libtpu.so: cannot open shared object file: No such file or directory"} {"num_params": 272938, "event": "Model created", "logger": "hw04", "level": "inf o", "timestamp": "2025-10-04T23:02:32.478968Z"} {"epoch": 1, "step": 704, "val_loss": 1.6888156540786163, "val_acc": 0.388844936 7088608, "event": "epoch_summary", "logger": "hw04.training", "level": "info", " timestamp": "2025-10-04T23:04:50.694858Z"} {"epoch": 2, "step": 1408, "val_loss": 1.317224457294126, "val_acc": 0.517207278 4810127, "event": "epoch_summary", "logger": "hw04.training", "level": "info", " timestamp": "2025-10-04T23:07:19.472678Z"} {"epoch": 3, "step": 2112, "val_loss": 1.2149084308479405, "val_acc": 0.55992879 74683544, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-04T23:09:49.454928Z"} {"epoch": 4, "step": 2816, "val_loss": 1.088192823566968, "val_acc": 0.612737341 7721519, "event": "epoch_summary", "logger": "hw04.training", "level": "info", " timestamp": "2025-10-04T23:43:07.676783Z"} {"event": "Finished blocking save. Time taken: 0.397554s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_5."} {"epoch": 5, "step": 3520, "val_loss": 0.9746726042107691, "val_acc": 0.65031645 56962026, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T00:24:08.390924Z"} {"epoch": 6, "step": 4224, "val_loss": 1.0020704873000519, "val_acc": 0.64477848 10126582, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T01:12:26.285167Z"} {"epoch": 7, "step": 4928, "val_loss": 0.8331410703779776, "val_acc": 0.71182753 16455697, "event": "epoch_summary", "logger": "hw04.training", "level": "info", </pre> | | |

| Oct 08, 2025 16:33 | train_cifar10.json | Page 2/4 |
|---|--------------------|----------|
| <pre> "timestamp": "2025-10-05T02:11:35.316912Z"} {"epoch": 8, "step": 5632, "val_loss": 0.8326036997988254, "val_acc": 0.70431170 88607594, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:23:58.950863Z"} {"epoch": 9, "step": 6336, "val_loss": 0.8269730994973001, "val_acc": 0.70530063 29113924, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:25:53.734331Z"} {"event": "Finished blocking save. Time taken: 0.291334s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_10."} {"epoch": 10, "step": 7040, "val_loss": 0.7053691990013364, "val_acc": 0.7565268 987341772, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:28:01.743996Z"} {"epoch": 11, "step": 7744, "val_loss": 0.7308083606671684, "val_acc": 0.7488132 911392406, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:30:21.243522Z"} {"epoch": 12, "step": 8448, "val_loss": 0.6420914360994026, "val_acc": 0.7759098 101265823, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:32:36.519162Z"} {"epoch": 13, "step": 9152, "val_loss": 0.6851127589050727, "val_acc": 0.7650316 455696202, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:35:02.441586Z"} {"epoch": 14, "step": 9856, "val_loss": 0.6394457692586923, "val_acc": 0.7806566 455696202, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:37:23.746143Z"} {"epoch": 15, "step": 10560, "val_loss": 0.7071164020254642, "val_acc": 0.756329 1139240507, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:39:50.385917Z"} {"event": "Finished async_save (blocking + background). Time taken: 0.507577s. d irectory=/work/hw04/checkpoints/state_epoch_15"} {"event": "[process=0][thread=async_save] Background save thread done. Time take n: 0.422445s."} {"epoch": 16, "step": 11264, "val_loss": 0.6119468215900131, "val_acc": 0.789952 5316455697, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:42:19.938507Z"} {"epoch": 17, "step": 11968, "val_loss": 0.5802336259733273, "val_acc": 0.801424 0506329114, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:45:14.132329Z"} {"epoch": 18, "step": 12672, "val_loss": 0.5539602588249158, "val_acc": 0.815466 7721518988, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:47:53.772006Z"} {"epoch": 19, "step": 13376, "val_loss": 0.5059745707843877, "val_acc": 0.822587 0253164557, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:50:45.324890Z"} {"event": "Finished blocking save. Time taken: 0.045058s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_20."} {"epoch": 20, "step": 14080, "val_loss": 0.5034029578106313, "val_acc": 0.824960 4430379747, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:53:26.954941Z"} {"event": "Finished async_save (blocking + background). Time taken: 0.561195s. d irectory=/work/hw04/checkpoints/best_state_epoch_20"} {"event": "[process=0][thread=async_save] Background save thread done. Time take n: 0.509544s."} {"epoch": 21, "step": 14784, "val_loss": 0.4975973575175563, "val_acc": 0.830893 9873417721, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:55:42.652465Z"} {"epoch": 22, "step": 15488, "val_loss": 0.5027272561682931, "val_acc": 0.830300 6329113924, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T02:57:59.114260Z"} {"epoch": 23, "step": 16192, "val_loss": 0.5082427591462678, "val_acc": 0.827136 0759493671, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T03:00:27.786364Z"} {"epoch": 24, "step": 16896, "val_loss": 0.4491530262593982, "val_acc": 0.842167 7215189873, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T03:02:49.746914Z"} {"event": "Finished blocking save. Time taken: 0.070533s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_25."} {"epoch": 25, "step": 17600, "val_loss": 0.45212958206104326, "val_acc": 0.84276 10759493671, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T03:05:19.481743Z"} </pre> | | |

| Oct 08, 2025 16:33 | train_cifar10.json | Page 3/4 |
|--|--------------------|----------|
| <pre>{ "event": "Finished async_save (blocking + background). Time taken: 0.538781s. d irectory=/work/hw04/checkpoints/best_state_epoch_25"} "event": "[process=0][thread=async_save] Background save thread done. Time take n: 0.454450s." "epoch": 26, "step": 18304, "val_loss": 0.44651460760756384, "val_acc": 0.84829 90506329114, "event": "epoch_summary", "logger": "hw04.training", "level": "info ", "timestamp": "2025-10-05T03:08:14.108608Z"} "epoch": 27, "step": 19008, "val_loss": 0.43895917282074315, "val_acc": 0.85027 68987341772, "event": "epoch_summary", "logger": "hw04.training", "level": "info ", "timestamp": "2025-10-05T03:10:33.586123Z"} "epoch": 28, "step": 19712, "val_loss": 0.4258923664500442, "val_acc": 0.857397 1518987342, "event": "epoch_summary", "logger": "hw04.training", "level": "info" , "timestamp": "2025-10-05T03:12:58.646980Z"} "epoch": 29, "step": 20416, "val_loss": 0.4334387503847291, "val_acc": 0.857594 9367088608, "event": "epoch_summary", "logger": "hw04.training", "level": "info" , "timestamp": "2025-10-05T03:15:29.196751Z"} "event": "Finished blocking save. Time taken: 0.114508s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_30." "epoch": 30, "step": 21120, "val_loss": 0.43960481146468394, "val_acc": 0.84988 1329113924, "event": "epoch_summary", "logger": "hw04.training", "level": "info" , "timestamp": "2025-10-05T03:18:00.676294Z"} "event": "Finished async_save (blocking + background). Time taken: 0.500118s. d irectory=/work/hw04/checkpoints/best_state_epoch_30"} "epoch": 31, "step": 21824, "val_loss": 0.41756087937687014, "val_acc": 0.85977 05696202531, "event": "epoch_summary", "logger": "hw04.training", "level": "info ", "timestamp": "2025-10-05T03:20:23.191371Z"} "epoch": 32, "step": 22528, "val_loss": 0.4153142190432247, "val_acc": 0.866297 4683544303, "event": "epoch_summary", "logger": "hw04.training", "level": "info" , "timestamp": "2025-10-05T03:22:51.648931Z"} "epoch": 33, "step": 23232, "val_loss": 0.41485248147686826, "val_acc": 0.86036 39240506329, "event": "epoch_summary", "logger": "hw04.training", "level": "info ", "timestamp": "2025-10-05T03:25:20.120943Z"} "epoch": 34, "step": 23936, "val_loss": 0.4226948241644268, "val_acc": 0.862737 3417721519, "event": "epoch_summary", "logger": "hw04.training", "level": "info" , "timestamp": "2025-10-05T03:28:03.197280Z"} "epoch": 35, "step": 24640, "val_loss": 0.41677922802635387, "val_acc": 0.86115 50632911392, "event": "epoch_summary", "logger": "hw04.training", "level": "info ", "timestamp": "2025-10-05T03:30:31.283716Z"} "event": "Finished async_save (blocking + background). Time taken: 0.437339s. d irectory=/work/hw04/checkpoints/best_state_epoch_35"} "epoch": 36, "step": 25344, "val_loss": 0.4113709775936127, "val_acc": 0.86056 17088607594, "event": "epoch_summary", "logger": "hw04.training", "level": "info ", "timestamp": "2025-10-05T03:33:27.668778Z"} "epoch": 37, "step": 26048, "val_loss": 0.40533167953732646, "val_acc": 0.86491 29746835443, "event": "epoch_summary", "logger": "hw04.training", "level": "info ", "timestamp": "2025-10-05T03:35:56.804861Z"} "epoch": 38, "step": 26752, "val_loss": 0.4060941059378129, "val_acc": 0.865110 7594936709, "event": "epoch_summary", "logger": "hw04.training", "level": "info" , "timestamp": "2025-10-05T03:38:24.808950Z"} "epoch": 39, "step": 27456, "val_loss": 0.4030414829148522, "val_acc": 0.867681 9620253164, "event": "epoch_summary", "logger": "hw04.training", "level": "info" , "timestamp": "2025-10-05T03:40:52.277778Z"} "event": "Finished blocking save. Time taken: 0.118553s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_40." "epoch": 40, "step": 28160, "val_loss": 0.404538533166994, "val_acc": 0.8651107 594936709, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-05T03:43:35.801138Z"} "event": "Training finished. Generating final plots...", "logger": "hw04.traini ng", "level": "info", "timestamp": "2025-10-05T03:43:35.804946Z"} "settings": { "debug": false, "random_seed": 31415, "model": { "input_depth": 3, "layer_depths": [16, 32, 64], "blocks_per_stage": [3, 3, 3], "num_classes": 10, "data": { "batch_size": 64, "validation_size": 5000, "shuffle_buffer": 10000, "cifar100": false, "training": { "num_epochs": 40, "learning_rate": 0.001, "weight_decay": 0.0001, "log_interval": 100, "final_test": true, "checkpoint_dir": "Pos ixPath('hw04/checkpoints_cifar10')", "plotting": { "output_dir": "PosixPath('hw0 4/artifacts')", "figsize": [5, 3], "dpi": 200 }, "event": "Settings loaded", "lo gger": "hw04", "level": "info", "timestamp": "2025-10-08T20:07:20.965938Z" } } }, "event": "Variant folder /root/tensorflow_datasets/cifar10/3.0.2 has no dataset _info.json" }</pre> | | |

| Oct 08, 2025 16:33 | train_cifar10.json | Page 4/4 |
|---|--------------------|----------|
| <pre>{ "event": "Generating dataset cifar10 (/root/tensorflow_datasets/cifar10/3.0.2)" } "event": "Downloading https://www.cs.toronto.edu/~kriz/cifar-10-binary.tar.gz i nto /root/tensorflow_datasets/downloads/cifar10/cs.toronto.edu_kriz_cifar-10-bin aryODHPtIjLh3oLcXirEIST07dkzyKjRCuo16lV8Wc6C7s.tar.gz.tmp.f89797db97604e0c9d9c28 774f134f32..." "event": "Done writing /root/tensorflow_datasets/cifar10/incomplete.DJLYYH_3.0. 2/cifar10-train.tfrecord*. Number of examples: 50000 (shards: [50000])" "event": "Done writing /root/tensorflow_datasets/cifar10/incomplete.DJLYYH_3.0. 2/cifar10-test.tfrecord*. Number of examples: 10000 (shards: [10000])" "event": "Creating a tf.data.Dataset reading 1 files located in folders: /root/ tensorflow_datasets/cifar10/3.0.2" "event": "Constructing tf.data.Dataset cifar10 for split train, from /root/tens orflow_datasets/cifar10/3.0.2" "event": "Load dataset info from /root/tensorflow_datasets/cifar10/3.0.2" "event": "Creating a tf.data.Dataset reading 1 files located in folders: /root/ tensorflow_datasets/cifar10/3.0.2" "event": "Constructing tf.data.Dataset cifar10 for split test, from /root/tenso rflow_datasets/cifar10/3.0.2" "event": "Unable to initialize backend 'tpu': INTERNAL: Failed to open libtpu.s o: libtpu.so: cannot open shared object file: No such file or directory" "num_params": 272938, "event": "Model created", "logger": "hw04", "level": "inf o", "timestamp": "2025-10-08T20:08:05.943964Z"} "event": "Starting final testing...", "logger": "hw04", "level": "info", "times tamp": "2025-10-08T20:08:06.201148Z"} "event": "orbax-checkpoint version: 0.11.25" "event": "Created BasePyTreeCheckpointHandler: use_ocdbt=True, use_zarr3=False, pytree_metadata_options=PyTreeMetadataOptions(support_rich_types=False), array_ metadata_store=<orbax.checkpoint._src.metadata.array_metadata_store.Store object at 0x7f2ffefd9b90>, enable_pinned_host_transfer=False, save_concurrent_bytes: 9 6000000000 (89.4 GiB), restore_concurrent_bytes: 96000000000 (89.4 GiB)" "event": "[thread=MainThread] Failed to get flag value for EXPERIMENTAL_ORBAX_U SE_DISTRIBUTED_PROCESS_ID." "event": "[process=0][thread=MainThread] Using barrier_sync_fn: <function get_b arrier_sync_fn.<locals>.<lambda> at 0x7f2f7a966160> timeout: 600 secs and primar y_host=0 for async checkpoint writes" "event": "Restoring checkpoint from /work/hw04/checkpoints_cifar10/best_state_e poch_40." "event": "[process=0] /jax/checkpoint/read/gbytes_per_sec: 0 Bytes/s (total gby tes: 1.0 MiB) (time elapsed: 296 milliseconds) (per-host)" "event": "Finished restoring checkpoint in 0.30 seconds from /work/hw04/checkpo ints_cifar10/best_state_epoch_40." "test_loss": 1.4637850788748188, "test_acc": 0.8576831210191083, "event": "test _results", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T 20:08:17.975697Z"} "test_loss": 1.4637850788748188, "test_acc": 0.8576831210191083, "event": "Fina l Test Results", "logger": "hw04", "level": "info", "timestamp": "2025-10-08T20: 08:18.006717Z"}</pre> | | |

Loss and Accuracy Graphs for CIFAR-10 Training



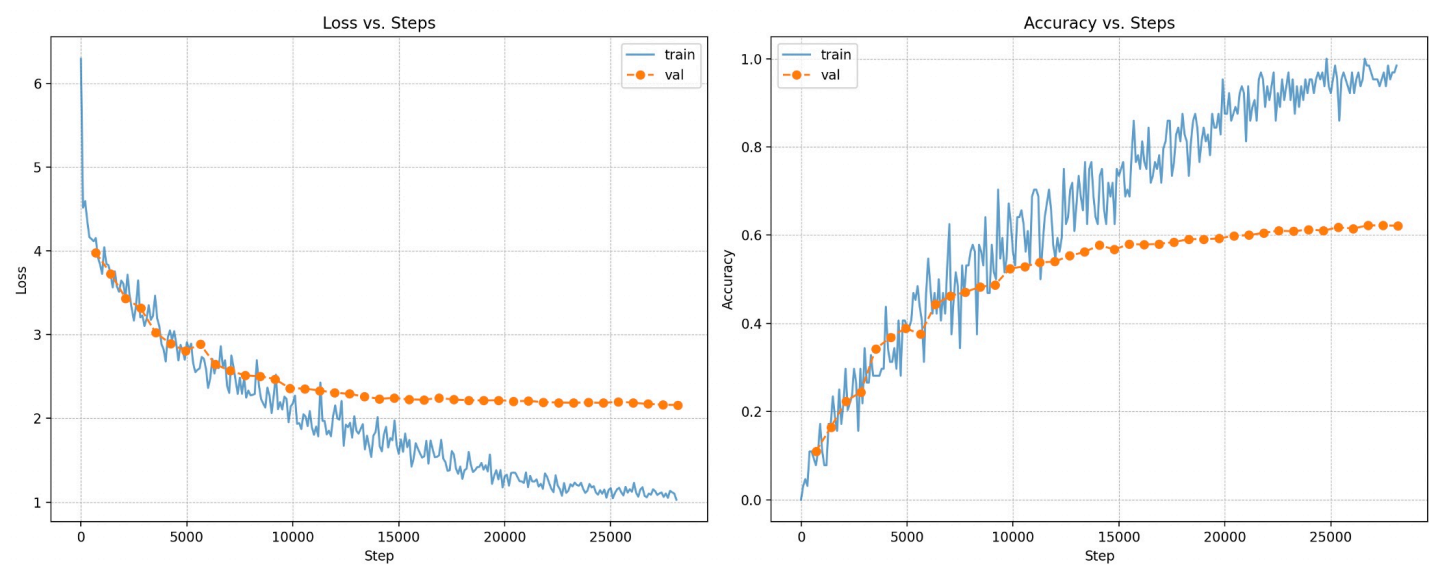
| Oct 08, 2025 16:42 | train_cifar100.json | Page 1/4 |
|--|---------------------|----------|
| <pre> /* I deleted all the training log (train loss) since there were too many and ju st leave the validation loss log It should be easier to see the train/validation loss graph to see the traini ng summary I see some overfitting from the validation loss graphs. Stronger augmentation and Dropout could be helpful Test set result is at the very bottom I tried to use renet-34 like structure but with less layer depths. Training took about 7 hours in my cpu. Started: 2025-10-08T02:04:23.076490Z Ended: 2025-10-08T09:18:38.831492Z */ {"settings": {"debug": false, "random_seed": 31415, "model": {"input_depth": 3, "layer_depths": [32, 64, 128, 256], "blocks_per_stage": [3, 4, 6, 3], "num_class es": 100}, "data": {"batch_size": 64, "validation_size": 5000, "shuffle_buffer": 10000, "cifar100": true}, "training": {"num_epochs": 40, "learning_rate": 0.001 , "weight_decay": 0.0001, "log_interval": 100, "final_test": false, "checkpoint_ dir": "PosixPath('hw04/checkpoints')"}, "plotting": {"output_dir": "PosixPath('h w04/artifacts')"}, "figsize": [5, 3], "dpi": 200}}, {"event": "Settings loaded", "logger": "hw04", "level": "info", "timestamp": "2025-10-08T02:03:14.897531Z"} {"event": "Variant folder /root/tensorflow_datasets/cifar100/3.0.2 has no datase t_info.json"} {"event": "Generating dataset cifar100 (/root/tensorflow_datasets/cifar100/3.0.2)"} {"event": "Downloading https://www.cs.toronto.edu/~kriz/cifar-100-binary.tar.gz into /root/tensorflow_datasets/downloads/cifar100/cs.toronto.edu-kriz_cifar-100- binaryzK0jb7CkNxmV4pH2clu5WdAIotsPlZhrMxx9-DELEk.tar.gz.tmp.05f7ache3c27433ab9d 9a7be715a262a..."} {"event": "Done writing /root/tensorflow_datasets/cifar100/incomplete.CJ1809_3.0 .2/cifar100-train.tfrecord*. Number of examples: 50000 (shards: [50000])"} {"event": "Done writing /root/tensorflow_datasets/cifar100/incomplete.CJ1809_3.0 .2/cifar100-test.tfrecord*. Number of examples: 10000 (shards: [10000])"} {"event": "Creating a tf.data.Dataset reading 1 files located in folders: /root/ tensorflow_datasets/cifar100/3.0.2."} {"event": "Constructing tf.data.Dataset cifar100 for split train, from /root/tens orflow_datasets/cifar100/3.0.2"} {"event": "Load dataset info from /root/tensorflow_datasets/cifar100/3.0.2"} {"event": "Creating a tf.data.Dataset reading 1 files located in folders: /root/ tensorflow_datasets/cifar100/3.0.2."} {"event": "Constructing tf.data.Dataset cifar100 for split test, from /root/tens orflow_datasets/cifar100/3.0.2"} {"event": "Unable to initialize backend 'tpu': INTERNAL: Failed to open libtpu.s o: libtpu.so: cannot open shared object file: No such file or directory"} {"num_params": 5352484, "event": "Model created", "logger": "hw04", "level": "in fo", "timestamp": "2025-10-08T02:04:23.076490Z"} {"epoch": 1, "step": 704, "val_loss": 3.976373268079154, "val_acc": 0.1097705696 2025317, "event": "epoch_summary", "logger": "hw04.training", "level": "info", " timestamp": "2025-10-08T02:12:48.870353Z"} {"epoch": 2, "step": 1408, "val_loss": 3.725985680954366, "val_acc": 0.164556962 02531644, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T02:23:02.027043Z"} {"epoch": 3, "step": 2112, "val_loss": 3.4336151533488986, "val_acc": 0.22329905 06329114, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T02:35:52.401048Z"} {"epoch": 4, "step": 2816, "val_loss": 3.315141044085539, "val_acc": 0.244066455 69620253, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T02:49:28.374315Z"} {"event": "Finished blocking save. Time taken: 0.405937s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_5."} {"epoch": 5, "step": 3520, "val_loss": 3.023668651339374, "val_acc": 0.341969936 7088608, "event": "epoch_summary", "logger": "hw04.training", "level": "info", " timestamp": "2025-10-08T03:03:20.608685Z"} {"epoch": 6, "step": 4224, "val_loss": 2.891840086707586, "val_acc": 0.368473101 2658228, "event": "epoch_summary", "logger": "hw04.training", "level": "info", " timestamp": "2025-10-08T03:15:15.281783Z"} </pre> | | |

| Oct 08, 2025 16:42 | train_cifar100.json | Page 2/4 |
|---|---------------------|----------|
| <pre> {"epoch": 7, "step": 4928, "val_loss": 2.80468063716647, "val_acc": 0.3890427215 1898733, "event": "epoch_summary", "logger": "hw04.training", "level": "info", " timestamp": "2025-10-08T03:30:54.135918Z"} {"epoch": 8, "step": 5632, "val_loss": 2.8846074689792682, "val_acc": 0.37539556 962025317, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T03:43:39.004092Z"} {"epoch": 9, "step": 6336, "val_loss": 2.644265703008145, "val_acc": 0.442840189 8734177, "event": "epoch_summary", "logger": "hw04.training", "level": "info", " timestamp": "2025-10-08T03:57:08.722001Z"} {"event": "Finished blocking save. Time taken: 0.287610s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_10."} {"epoch": 10, "step": 7040, "val_loss": 2.5681378705592097, "val_acc": 0.4622231 012658228, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T04:11:06.506207Z"} {"epoch": 11, "step": 7744, "val_loss": 2.513141556631161, "val_acc": 0.47072784 810126583, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T04:24:44.379628Z"} {"epoch": 12, "step": 8448, "val_loss": 2.5028848678250855, "val_acc": 0.4827927 2151898733, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T04:37:42.487969Z"} {"epoch": 13, "step": 9152, "val_loss": 2.4689806473406057, "val_acc": 0.4869462 0253164556, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T04:51:23.553805Z"} {"epoch": 14, "step": 9856, "val_loss": 2.3596031695981567, "val_acc": 0.5245253 164556962, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T05:05:20.270818Z"} {"event": "Finished blocking save. Time taken: 0.353670s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_15."} {"epoch": 15, "step": 10560, "val_loss": 2.3550007765806176, "val_acc": 0.528876 582278481, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T05:18:59.185930Z"} {"epoch": 16, "step": 11264, "val_loss": 2.332110721853715, "val_acc": 0.5379746 835443038, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T05:31:56.445378Z"} {"epoch": 17, "step": 11968, "val_loss": 2.3060759547390517, "val_acc": 0.539952 5316455697, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T05:42:01.686646Z"} {"epoch": 18, "step": 12000, "loss": 2.1541953086853027, "acc": 0.546875, "event ": "train_iter", "logger": "hw04.training", "level": "info", "timestamp": "2025- 10-08T05:42:26.991727Z"} {"epoch": 18, "step": 12100, "loss": 1.9982845783233643, "acc": 0.59375, "event" ": "train_iter", "logger": "hw04.training", "level": "info", "timestamp": "2025-1 0-08T05:43:40.246351Z"} {"epoch": 18, "step": 12200, "loss": 1.9845942258834839, "acc": 0.5625, "event" ": "train_iter", "logger": "hw04.training", "level": "info", "timestamp": "2025-1 0-08T05:44:54.006460Z"} {"epoch": 18, "step": 12300, "loss": 2.2072534561157227, "acc": 0.59375, "event" ": "train_iter", "logger": "hw04.training", "level": "info", "timestamp": "2025-1 0-08T05:46:06.017281Z"} {"epoch": 18, "step": 12400, "loss": 1.6721937656402588, "acc": 0.75, "event": " train_iter", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-0 8T05:47:17.964478Z"} {"epoch": 18, "step": 12500, "loss": 1.924797773361206, "acc": 0.625, "event": " train_iter", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-0 8T05:48:27.335253Z"} {"epoch": 18, "step": 12600, "loss": 1.895101547241211, "acc": 0.640625, "event" ": "train_iter", "logger": "hw04.training", "level": "info", "timestamp": "2025-1 0-08T05:49:37.582857Z"} {"epoch": 18, "step": 12672, "val_loss": 2.2927366072618507, "val_acc": 0.553006 329113924, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T05:50:42.273608Z"} {"epoch": 19, "step": 13376, "val_loss": 2.25943695442586, "val_acc": 0.56190664 55696202, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T06:02:30.057835Z"} {"event": "Finished blocking save. Time taken: 0.302445s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_20."} {"epoch": 20, "step": 14080, "val_loss": 2.2340246137184434, "val_acc": 0.577136 0759493671, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T06:13:47.481533Z"} </pre> | | |

| Oct 08, 2025 16:42 | train_cifar100.json | Page 3/4 |
|--|---------------------|----------|
| <pre>{"epoch": 21, "step": 14784, "val_loss": 2.243983276282685, "val_acc": 0.5674446202531646, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T06:23:34.818738Z"} {"epoch": 22, "step": 15488, "val_loss": 2.2258135306684275, "val_acc": 0.5797072784810127, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T06:32:59.383679Z"} {"epoch": 23, "step": 16192, "val_loss": 2.222693950315065, "val_acc": 0.5787183544303798, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T06:42:14.872808Z"} {"epoch": 24, "step": 16896, "val_loss": 2.243938494332229, "val_acc": 0.5801028481012658, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T06:52:52.934472Z"} {"event": "Finished blocking save. Time taken: 0.174913s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_25."} {"epoch": 25, "step": 17600, "val_loss": 2.224066859559168, "val_acc": 0.584256329113924, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T07:02:52.741474Z"} {"epoch": 26, "step": 18304, "val_loss": 2.2172716629656057, "val_acc": 0.5909810126582279, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T07:13:10.497284Z"} {"epoch": 27, "step": 19008, "val_loss": 2.2137932611417166, "val_acc": 0.5909810126582279, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T07:22:51.116493Z"} {"epoch": 28, "step": 19712, "val_loss": 2.216762476329562, "val_acc": 0.5923655063291139, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T07:32:24.958280Z"} {"epoch": 29, "step": 20416, "val_loss": 2.2051513798629183, "val_acc": 0.5986946202531646, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T07:42:18.802216Z"} {"epoch": 30, "step": 21120, "val_loss": 2.2079105467735967, "val_acc": 0.599881329113924, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T07:52:27.256982Z"} {"event": "Finished async_save (blocking + background). Time taken: 1.108641s. directory=/work/hw04/checkpoints/best_state_epoch_30"} {"event": "[process=0][thread=async_save] Background save thread done. Time taken: 0.928715s."} {"epoch": 31, "step": 21824, "val_loss": 2.1937194051621836, "val_acc": 0.6050237341772152, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T08:02:25.895055Z"} {"epoch": 32, "step": 22528, "val_loss": 2.1885239006597783, "val_acc": 0.6099683544303798, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T08:12:54.682985Z"} {"epoch": 33, "step": 23232, "val_loss": 2.186336153670202, "val_acc": 0.6087816455696202, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T08:23:14.888357Z"} {"epoch": 34, "step": 23936, "val_loss": 2.1903039745137662, "val_acc": 0.6125395569620253, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T08:31:10.139681Z"} {"event": "Finished blocking save. Time taken: 0.171701s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_35."} {"epoch": 35, "step": 24640, "val_loss": 2.183525858046133, "val_acc": 0.6101661392405063, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T08:39:03.468770Z"} {"epoch": 36, "step": 25344, "val_loss": 2.197526878948453, "val_acc": 0.6176819620253164, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T08:46:58.679913Z"} {"epoch": 37, "step": 26048, "val_loss": 2.1875499139858197, "val_acc": 0.6149129746835443, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T08:54:52.045278Z"} {"epoch": 38, "step": 26752, "val_loss": 2.171840216540083, "val_acc": 0.6220332278481012, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T09:02:46.707634Z"} {"epoch": 39, "step": 27456, "val_loss": 2.1635287939747676, "val_acc": 0.6218354430379747, "event": "epoch_summary", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T09:10:44.956816Z"} {"event": "Finished blocking save. Time taken: 0.342879s. Continuing background save to /work/hw04/checkpoints/best_state_epoch_40."} {"epoch": 40, "step": 28160, "val_loss": 2.158977393862567, "val_acc": 0.6214398734177216, "event": "epoch_summary", "logger": "hw04.training", "level": "info",</pre> | | |

| Oct 08, 2025 16:42 | train_cifar100.json | Page 4/4 |
|---|---------------------|----------|
| <pre>"timestamp": "2025-10-08T09:18:38.831492Z"} {"num_params": 5352484, "event": "Model created", "logger": "hw04", "level": "info", "timestamp": "2025-10-08T19:58:42.569048Z"} {"event": "Starting final testing...", "logger": "hw04", "level": "info", "timestamp": "2025-10-08T19:58:42.898115Z"} {"event": "orbax-checkpoint version: 0.11.25"} {"event": "Created BasePyTreeCheckpointHandler: use_ocdbt=True, use_zarr3=False, pytree_metadata_options=PyTreeMetadataOptions(support_rich_types=False), array_metadata_store=<orbax.checkpoint._src.metadata.array_metadata_store.Store object at 0x7f589abd6790>, enable_pinned_host_transfer=False, save_concurrent_bytes: 96000000000 (89.4 GiB), restore_concurrent_bytes: 96000000000 (89.4 GiB)"} {"event": "[thread=MainThread] Failed to get flag value for EXPERIMENTAL_ORBAX_USE_DISTRIBUTED_PROCESS_ID."} {"event": "[process=0][thread=MainThread] Using barrier_sync_fn: <function get_barrier_sync_fn.<locals>.<lambda> at 0x7f5816f520c0> timeout: 600 secs and primary_host=0 for async checkpoint writes"} {"event": "Restoring checkpoint from /work/hw04/checkpoints/best_state_epoch_40."} {"event": "[process=0] /jax/checkpoint/read/gbytes_per_sec: 0 Bytes/s (total gbytes: 20.4 MiB) (time elapsed: 549 milliseconds) (per-host)"} {"event": "Finished restoring checkpoint in 0.56 seconds from /work/hw04/checkpoints/best_state_epoch_40."} {"test_loss": 2.175596616830036, "test_acc": 0.6136544585987261, "event": "test_results", "logger": "hw04.training", "level": "info", "timestamp": "2025-10-08T19:59:12.667384Z"} {"test_loss": 2.175596616830036, "test_acc": 0.6136544585987261, "event": "Final Test Results", "logger": "hw04", "level": "info", "timestamp": "2025-10-08T19:59:12.705750Z"}</pre> | | |

Loss and Accuracy Graphs for CIFAR-100 Training



Sep 28, 2025 17:36

pyproject.toml

Page 1/1

```
# pyproject.toml.jinja

[project]
name = "hw04"
version = "0.1.0"
description = "CIFARIO"
readme = "README.md"
authors = [
    { name = "Wongee (Freddy) Hong", email = "wongee.hong@cooper.edu" }
]
requires-python = ">=3.11"
dependencies = [
    "structlog",
    "numpy",
    "tensorflow",
    "tensorflow_datasets",
    "pydantic-settings",
    "matplotlib",
    "tqdm",
    "jax",
    "jaxlib",
    "flax",
    "optax",
    "scikit-learn",
]

[project.scripts]
hw04 = "hw04:main"

[build-system]
requires = ["uv_build>=0.8.3,<0.9.0"]
build-backend = "uv_build"
```

Oct 04, 2025 17:22

logging.py

Page 1/2

```

import logging
import os
import sys
from pathlib import Path

import jax
import numpy as np
import structlog

class FormattedFloat(float):
    def __repr__(self) -> str:
        return f"{self:4g}"

def custom_serializer_processor(logger, method_name, event_dict):
    for key, value in event_dict.items():
        if hasattr(value, "numpy"):
            value = value.numpy()
        if isinstance(value, jax.Array):
            value = np.array(value)
        if isinstance(value, (np.generic, np.ndarray)):
            value = value.item() if value.size == 1 else value.tolist()
        if isinstance(value, float):
            value = FormattedFloat(value)
        if isinstance(value, Path):
            value = str(value)
        event_dict[key] = value
    return event_dict

def configure_logging(log_dir: Path = Path("hw04/artifacts"), log_name: str = "train.js
n"):
    log_dir.mkdir(parents=True, exist_ok=True)
    log_file = log_dir / log_name
    log_level = os.environ.get("LOG_LEVEL", "INFO").upper()

    shared_processors = [
        structlog.stdlib.add_logger_name,
        structlog.stdlib.add_log_level,
        structlog.stdlib.PositionalArgumentsFormatter(),
        structlog.processors.TimeStamper(fmt="iso"),
        structlog.processors.StackInfoRenderer(),
        structlog.processors.format_exc_info,
        structlog.processors.UnicodeDecoder(),
        custom_serializer_processor,
    ]

    structlog.configure(
        processors=shared_processors + [

            structlog.stdlib.ProcessorFormatter.wrap_for_formatter,
        ],
        logger_factory=structlog.stdlib.LoggerFactory(),
        wrapper_class=structlog.stdlib.BoundLogger,
        cache_logger_on_first_use=True,
    )

    console_renderer = structlog.dev.ConsoleRenderer(
        colors=sys.stdout.isatty(),
        exception_formatter=structlog.dev.RichTracebackFormatter(),
    )
    console_handler = logging.StreamHandler(sys.stdout)
    console_handler.setFormatter(
        structlog.stdlib.ProcessorFormatter(

            processors=[
                structlog.stdlib.ProcessorFormatter.remove_processors_meta,
                console_renderer

```

Oct 04, 2025 17:22

logging.py

Page 2/2

```

        ],
    )

    file_renderer = structlog.processors.JSONRenderer()
    file_handler = logging.FileHandler(log_file, mode="w")
    file_handler.setFormatter(
        structlog.stdlib.ProcessorFormatter(
            processors=[
                structlog.stdlib.ProcessorFormatter.remove_processors_meta,
                file_renderer
            ],
        )
    )

    root_logger = logging.getLogger()
    root_logger.addHandler(console_handler)
    root_logger.addHandler(file_handler)
    root_logger.setLevel(log_level)

    logging.getLogger("matplotlib").setLevel(logging.WARNING)
    logging.getLogger("PIL").setLevel(logging.WARNING)

    return log_file

```


Oct 08, 2025 16:42

config.py

Page 1/1

```

from pathlib import Path
from typing import Tuple
from pydantic import BaseModel
from pydantic_settings import BaseSettings

class ModelSettings(BaseModel):
    input_depth: int = 3
    layer_depths: list[int] = [32, 64, 128, 256]
    blocks_per_stage: list[int] = [3, 4, 6, 3]
    num_classes: int = 100

class DataSettings(BaseModel):
    batch_size: int = 64
    validation_size: int = 5000
    shuffle_buffer: int = 10000
    cifar100: bool = True

class TrainingSettings(BaseModel):
    """Settings for model training."""
    num_epochs: int = 40
    learning_rate: float = 0.001
    weight_decay: float = 1e-4          # L2 penalty
    log_interval: int = 100
    final_test: bool = True
    checkpoint_dir: Path = Path("hw04/checkpoints")

class PlottingSettings(BaseModel):
    """Settings for logging and saving artifacts."""
    output_dir: Path = Path("hw04/artifacts")
    figsize: Tuple[int, int] = (5, 3)
    dpi: int = 200

class AppSettings(BaseSettings):
    """Main application settings."""
    debug: bool = False
    random_seed: int = 31415

    model: ModelSettings = ModelSettings()
    data: DataSettings = DataSettings()
    training: TrainingSettings = TrainingSettings()
    plotting: PlottingSettings = PlottingSettings()

def load_settings() -> AppSettings:
    """Load application settings."""
    return AppSettings()

```

Oct 08, 2025 16:47

__init__.py

Page 1/2

```

import optax
import structlog
from flax import nnx
import jax
import pathlib as Path
import orbax.checkpoint as ocp

from .config import load_settings
from .logging import configure_logging
from .model import Classifier
from .training import train, test_evaluation, load_checkpoint
from .data import load_CIFAR10, load_CIFAR100

def main() -> None:
    settings = load_settings()
    log_file = configure_logging()
    log = structlog.get_logger()
    log.info("Settings loaded", settings=settings.model_dump())
    print(f"Logs are being saved to {log_file}")

    if settings.data.cifar100:
        ds_train, ds_val, ds_test = load_CIFAR100(
            batch_size=settings.data.batch_size,
            validation_size=settings.data.validation_size,
        )
    else:
        ds_train, ds_val, ds_test = load_CIFAR10(
            batch_size=settings.data.batch_size,
            validation_size=settings.data.validation_size,
        )
    rngs = nnx.Rngs(params=settings.random_seed)
    model = Classifier(
        rngs=rngs,
        input_depth=settings.model.input_depth,
        layer_depths=settings.model.layer_depths,
        blocks_per_stage=settings.model.blocks_per_stage,
        layer_kernel_sizes=[(3, 3)] * len(settings.model.layer_depths),
        num_classes=settings.model.num_classes,
    )

    params = nnx.state(model, nnx.Param)
    num_params = sum(p.size for p in jax.tree_util.tree_leaves(params))
    log.info("Model created", num_params=num_params)

    # LR schedule and optimizer
    train_size = 50000
    steps_per_epoch = train_size // settings.data.batch_size
    decay_steps = settings.training.num_epochs * steps_per_epoch
    lr_scheduler = optax.schedules.cosine_decay_schedule(
        init_value=settings.training.learning_rate,
        decay_steps=decay_steps,
    )
    optimizer = nnx.Optimizer(
        model,
        optax.adamw(
            learning_rate=lr_scheduler,
            weight_decay=settings.training.weight_decay,
        ),
        wrt=nnx.Param,
    )

    if settings.training.final_test:
        log.info("Starting final testing...")

        model_kwargs = dict(
            input_depth=settings.model.input_depth,
            layer_depths=settings.model.layer_depths,
            blocks_per_stage=settings.model.blocks_per_stage,
            layer_kernel_sizes=[(3, 3)] * len(settings.model.layer_depths),

```

Oct 08, 2025 16:47

__init__.py

Page 2/2

```

        num_classes=settings.model.num_classes,
    )

    model = load_checkpoint(
        Classifier,
        model_kwargs
    )

    test_loss, test_acc = test_evaluation(model, ds_test, rngs)
    log.info("Final Test Results", test_loss=test_loss, test_acc=test_acc)
else:
    _ = train(model, optimizer, ds_train, ds_val, rngs)

```

Oct 08, 2025 16:47

model.py

Page 1/3

Deleted dropout to implement checkpoint

```
import jax
import jax.numpy as jnp
from flax import nnx
```

```
class Conv2d(nnx.Module):
```

```
    def __init__(
        self,
        in_channels: int,
        out_channels: int,
        kernel_size,
        *,
        strides=(1, 1),
        padding="SAME",
        use_bias=True,
        rngs: nnx.Rngs,
    ):
        self.conv = nnx.Conv(
            in_channels, out_channels, kernel_size,
            strides=strides, padding=padding, use_bias=use_bias, rngs=rngs
        )
```

```
    def __call__(self, x, *, rngs: nnx.Rngs):
        x = self.conv(x)
        return x
```

```
class ResidualBlock(nnx.Module):
```

```
    """Basic residual block."""
```

```
    def __init__(
        self,
        in_channels: int,
        out_channels: int,
        *,
        stride=(1, 1),
        rngs: nnx.Rngs,
    ):
        self.norm1 = nnx.GroupNorm(num_features=in_channels, num_groups=8, rngs=rngs)
        self.conv1 = Conv2d(
            in_channels, out_channels, (3, 3),
            strides=stride, rngs=rngs,
        )
        self.norm2 = nnx.GroupNorm(num_features=out_channels, num_groups=8, rngs=rngs)
        self.conv2 = Conv2d(
            out_channels, out_channels, (3, 3),
            strides=(1, 1), rngs=rngs,
        )
        self.shortcut = (
            Conv2d(
                in_channels, out_channels, (1, 1),
                strides=stride, rngs=rngs,
            )
            if in_channels != out_channels or stride != (1, 1)
            else None
        )
```

```
    def __call__(self, x, *, rngs: nnx.Rngs):
        out = self.norm1(x)
        out = nnx.relu(out)
        out = self.conv1(out, rngs=rngs)

        out = self.norm2(out)
        out = nnx.relu(out)
        out = self.conv2(out, rngs=rngs)
```

Oct 08, 2025 16:47

model.py

Page 2/3

```
identity = x if self.shortcut is None else self.shortcut(x, rngs=rngs)
return out + identity
```

```
class ResidualStage(nnx.Module):
```

```
    def __init__(
        self,
        in_channels: int,
        out_channels: int,
        num_blocks: int,
        stride=(1, 1),
        *,
        rngs: nnx.Rngs,
    ):
        blocks = []
        blocks.append(
            ResidualBlock(
                in_channels,
                out_channels,
                stride=stride,
                rngs=rngs,
            )
        )
        # Rest keep stride = (1,1)
        for _ in range(1, num_blocks):
            blocks.append(
                ResidualBlock(
                    out_channels,
                    out_channels,
                    stride=(1, 1),
                    rngs=rngs,
                )
            )
        self.blocks = nnx.List(blocks)
```

```
    def __call__(self, x, *, rngs: nnx.Rngs):
        for block in self.blocks:
            x = block(x, rngs=rngs)
        return x
```

```
class Classifier(nnx.Module):
```

```
    """Residual network classifier for MNIST."""
```

```
    def __init__(
        self,
        input_depth: int,
        layer_depths: list[int],
        blocks_per_stage,
        layer_kernel_sizes: list[tuple[int, int]],
        num_classes: int,
        *,
        rngs: nnx.Rngs,
    ):
        assert len(layer_depths) == len(layer_kernel_sizes), \
            "layer_depths and layer_kernel_sizes must match in length"
        self.init_conv = Conv2d(input_depth, layer_depths[0],
                                kernel_size=layer_kernel_sizes[0],
                                strides=(1, 1), rngs=rngs)

        # Build stages
        stages = []
        self.blocks_per_stage = blocks_per_stage
        in_ch = layer_depths[0]
        for i, out_ch in enumerate(layer_depths):
            stride = (2, 2) if i > 0 else (1, 1) # downsample at stage boundaries
            num_blocks = (self.blocks_per_stage[i]
                           if self.blocks_per_stage is not None
                           else 1)
            stage = ResidualStage(
```

```
        in_channels=in_ch,
        out_channels=out_ch,
        num_blocks=num_blocks,
        stride=stride,
        rngs=rngs,
    )
    stages.append(stage)
    in_ch = out_ch

self.stages = nnx.List(stages)

# Final classifier head
self.fc = nnx.Linear(in_ch, num_classes, rngs=rngs)

def __call__(self, x, *, rngs: nnx.Rngs):
    x = self.init_conv(x, rngs=rngs)
    for stage in self.stages:
        x = stage(x, rngs=rngs)

    x = jnp.mean(x, axis=(1, 2))

    # Final logits
    return self.fc(x)
```

Oct 08, 2025 16:46

training.py

Page 1/3

```

import jax
import jax.numpy as jnp
import optax
from flax import nnx
import matplotlib.pyplot as plt
import structlog
import orbax.checkpoint as ocp
from pathlib import Path
import os

from .config import load_settings

log = structlog.get_logger()

def cross_entropy_loss(logits, labels, num_classes, smoothing = 0.1):
    one_hot = jax.nn.one_hot(labels, num_classes)
    smoothed = one_hot * (1 - smoothing) + smoothing / num_classes # Implemented
    label smoothing for CIFAR-100 training.
    loss = optax.softmax_cross_entropy(logits, smoothed).mean()
    return loss

def compute_accuracy(logits, labels):
    preds = jnp.argmax(logits, axis=-1)
    return jnp.mean(preds == labels)

def _loss_fn(model, batch, rngs):
    settings = load_settings()
    images, labels = batch
    labels = labels.ravel()
    logits = model(images, rngs=rngs)
    loss = cross_entropy_loss(logits, labels, settings.model.num_classes)
    return loss, logits

@nnx.jit
def train_step(model, optimizer, batch, rngs):
    grad_fn = nnx.value_and_grad(_loss_fn, has_aux=True)
    (loss, logits), grads = grad_fn(model, batch, rngs)
    optimizer.update(model, grads)
    return loss, logits

@nnx.jit
def eval_step(model, batch, rngs):
    loss, logits = _loss_fn(model, batch, rngs)
    return loss, logits

def evaluate(model, dataset, rngs):
    total_loss, total_acc, n_batches = 0.0, 0.0, 0
    for batch in dataset.as_numpy_iterator():
        loss, logits = eval_step(model, batch, rngs)
        _, labels = batch
        labels = labels.ravel()
        acc = compute_accuracy(logits, labels)
        loss = float(loss)
        acc = float(acc)
        total_loss += loss
        total_acc += float(acc)
        n_batches += 1
    return total_loss / max(n_batches, 1), total_acc / max(n_batches, 1)

def train(model, optimizer, train_ds, val_ds, rngs):
    settings = load_settings()
    output_dir = settings.plotting.output_dir
    output_dir.mkdir(parents=True, exist_ok=True)
    checkpoint_dir = settings.training.checkpoint_dir.resolve()
    checkpoint_dir.mkdir(parents=True, exist_ok=True)

```

Oct 08, 2025 16:46

training.py

Page 2/3

```

checkpoint_dir = ocp.test_utils.erase_and_create_empty(checkpoint_dir)

log_interval = settings.training.log_interval
num_epochs = settings.training.num_epochs

metrics_history = {
    "train_steps": [],
    "train_loss": [],
    "train_accuracy": [],
    "val_steps": [],
    "val_loss": [],
    "val_accuracy": [],
}

global_step = 0
best_val_acc = 0.0
best_ckpt_path = None

for epoch in range(1, num_epochs + 1):
    for batch in train_ds.as_numpy_iterator():
        loss, logits = train_step(model, optimizer, batch, rngs)

        if global_step % log_interval == 0:
            images, labels = batch
            labels = labels.ravel()
            acc = compute_accuracy(logits, labels)

            loss_val = float(loss)
            acc_val = float(acc)

            metrics_history["train_steps"].append(global_step)
            metrics_history["train_loss"].append(loss_val)
            metrics_history["train_accuracy"].append(acc_val)

            log.info(
                "train_iter",
                epoch=epoch,
                step=global_step,
                loss=loss_val,
                acc=acc_val,
            )
            global_step += 1

    # Create checkpoint every 5 epoch. Compare with previous best validation
    # accuracy and save best checkpoint.
    val_loss, val_acc = evaluate(model, val_ds, rngs)
    if epoch % 5 == 0 or epoch == num_epochs:
        checkpointer = ocp.StandardCheckpointer()
        if val_acc > best_val_acc:
            best_val_acc = val_acc
            best_ckpt_path = checkpoint_dir / f"best_state_epoch_{epoch}"
            _, state = nnx.split(model)
            checkpointer.save(best_ckpt_path, state)
        else:
            _, state = nnx.split(model)
            ckpt_path = checkpoint_dir / f"state_epoch_{epoch}"
            checkpointer.save(ckpt_path, state)

    metrics_history["val_steps"].append(global_step)
    metrics_history["val_loss"].append(val_loss)
    metrics_history["val_accuracy"].append(val_acc)

    log.info(
        "epoch_summary",
        epoch=epoch,
        step=global_step,
        val_loss=val_loss,
        val_acc=val_acc,
    )

```

Oct 08, 2025 16:46

training.py

Page 3/3

```

    )

    log.info("Training finished. Generating final plots...")
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))

    ax1.set_title("Loss vs. Steps")
    ax2.set_title("Accuracy vs. Steps")

    ax1.plot(metrics_history["train_steps"], metrics_history["train_loss"], label="train",
, alpha=0.7)
    ax2.plot(metrics_history["train_steps"], metrics_history["train_accuracy"], label="train",
, alpha=0.7)

    ax1.plot(metrics_history["val_steps"], metrics_history["val_loss"], label="val", marker='o', linestyle='--')
    ax2.plot(metrics_history["val_steps"], metrics_history["val_accuracy"], label="val", marker='o', linestyle='--')

    ax1.set_xlabel("Step")
    ax1.set_ylabel("Loss")
    ax1.legend()
    ax1.grid(True, which='both', linestyle='--', linewidth=0.5)

    ax2.set_xlabel("Step")
    ax2.set_ylabel("Accuracy")
    ax2.legend()
    ax2.grid(True, which='both', linestyle='--', linewidth=0.5)

    fig.tight_layout()
    fig.savefig(output_dir / "final_training_metrics.png", dpi=settings.plotting.dpi)
    plt.close(fig)

    return metrics_history

def test_evaluation(model, test_ds, rngs):
    settings = load_settings()
    output_dir = settings.plotting.output_dir
    output_dir.mkdir(parents=True, exist_ok=True)

    test_loss, test_acc = evaluate(model, test_ds, rngs)
    log.info("test_results", test_loss=test_loss, test_acc=test_acc)

    return test_loss, test_acc

def load_checkpoint(model_cls, model_kwargs: dict):
    settings = load_settings()
    checkpoint_dir = settings.training.checkpoint_dir.resolve()
    checkpoint = ocp.StandardCheckpoint()

    ckpts = sorted(checkpoint_dir.glob("best_state_epoch_*"), key=os.path.getmtime)
    if not ckpts:
        raise FileNotFoundError(f"No checkpoints found in {checkpoint_dir}")

    latest_ckpt = ckpts[-1]
    print(f>Loading checkpoint from: {latest_ckpt}")

    abstract_model = nnx.eval_shape(lambda: model_cls(**model_kwargs, rngs=nnx.RNGs(0)))
    graphdef, abstract_state = nnx.split(abstract_model)

    # restore parameters
    state_restored = checkpoint.restore(latest_ckpt, abstract_state)
    model = nnx.merge(graphdef, state_restored)
    print("Checkpoint successfully restored!")
    return model

```

Oct 05, 2025 21:56

data.py

Page 1/2

```

import tensorflow_datasets as tfds
import tensorflow as tf

def load_CIFAR10(batch_size=128, validation_size=5000, seed=0):
    ds_train = tfds.load("cifar10", split="train", shuffle_files=True, as_supervised=True)
    ds_test = tfds.load("cifar10", split="test", as_supervised=True)

    def _preprocess(image, label):
        image = tf.cast(image, tf.float32) / 255.0
        label = tf.cast(label, tf.int32)
        return image, label

    ds_train = ds_train.map(_preprocess, num_parallel_calls=tf.data.AUTOTUNE)
    ds_test = ds_test.map(_preprocess, num_parallel_calls=tf.data.AUTOTUNE)

    ds_val = ds_train.take(validation_size)
    ds_train = ds_train.skip(validation_size)

    def augmentation(image, label):
        image = tf.image.resize_with_crop_or_pad(image, 40, 40)
        image = tf.image.random_crop(image, size=[32, 32, 3])
        image = tf.image.random_flip_left_right(image)
        image = tf.image.random_brightness(image, max_delta=0.2)
        image = tf.image.random_contrast(image, 0.8, 1.2)
        image = tf.image.random_saturation(image, 0.8, 1.2)
        return image, label

    ds_train = ds_train.map(augmentation, num_parallel_calls=tf.data.AUTOTUNE)

    ds_train = ds_train.shuffle(10 * batch_size, seed=seed).batch(batch_size).prefetch(
        tf.data.AUTOTUNE)
    ds_val = ds_val.batch(batch_size).prefetch(tf.data.AUTOTUNE)
    ds_test = ds_test.batch(batch_size).prefetch(tf.data.AUTOTUNE)

    return ds_train, ds_val, ds_test

def load_CIFAR100(batch_size=128, validation_size=5000, seed=0):
    ds_train = tfds.load("cifar100", split="train", shuffle_files=True, as_supervised=True)
    ds_test = tfds.load("cifar100", split="test", as_supervised=True)

    def _preprocess(image, label):
        image = tf.cast(image, tf.float32) / 255.0
        label = tf.cast(label, tf.int32)
        return image, label

    ds_train = ds_train.map(_preprocess, num_parallel_calls=tf.data.AUTOTUNE)
    ds_test = ds_test.map(_preprocess, num_parallel_calls=tf.data.AUTOTUNE)

    ds_val = ds_train.take(validation_size)
    ds_train = ds_train.skip(validation_size)

    def augmentation(image, label):
        image = tf.image.resize_with_crop_or_pad(image, 40, 40)
        image = tf.image.random_crop(image, size=[32, 32, 3])
        image = tf.image.random_flip_left_right(image)
        image = tf.image.random_brightness(image, max_delta=0.2)
        image = tf.image.random_contrast(image, 0.8, 1.2)
        image = tf.image.random_saturation(image, 0.8, 1.2)
        return image, label

    ds_train = ds_train.map(augmentation, num_parallel_calls=tf.data.AUTOTUNE)

    ds_train = ds_train.shuffle(10 * batch_size, seed=seed).batch(batch_size).prefetch(
        tf.data.AUTOTUNE)
    ds_val = ds_val.batch(batch_size).prefetch(tf.data.AUTOTUNE)
    ds_test = ds_test.batch(batch_size).prefetch(tf.data.AUTOTUNE)

```

Oct 05, 2025 21:56

data.py

Page 2/2

```

return ds_train, ds_val, ds_test

```