

Sep 22, 2025 17:48

pyproject.toml

Page 1/1

```
# pyproject.toml.jinja

[project]
name = "hw03"
version = "0.1.0"
description = "MNIST CNN"
readme = "README.md"
authors = [
    { name = "Wongee (Freddy) Hong", email = "wongee.hong@cooper.edu" }
]
requires-python = ">=3.11"
dependencies = [
    "structlog",
    "numpy",
    "tensorflow",
    "tensorflow_datasets",
    "pydantic-settings",
    "matplotlib",
    "tqdm",
    "jax",
    "jaxlib",
    "flax",
    "optax",
    "scikit-learn",
]

[project.scripts]
hw03 = "hw03:main"

[build-system]
requires = ["uv_build>=0.8.3,<0.9.0"]
build-backend = "uv_build"
```

Sep 24, 2025 21:57

logging.py

Page 1/2

```

import logging
import os
import sys
from pathlib import Path

import jax
import numpy as np
import structlog

class FormattedFloat(float):
    def __repr__(self) -> str:
        return f"{self:4g}"

def custom_serializer_processor(logger, method_name, event_dict):
    for key, value in event_dict.items():
        if hasattr(value, "numpy"):
            value = value.numpy()
        if isinstance(value, jax.Array):
            value = np.array(value)
        if isinstance(value, (np.generic, np.ndarray)):
            value = value.item() if value.size == 1 else value.tolist()
        if isinstance(value, float):
            value = FormattedFloat(value)
        if isinstance(value, Path):
            value = str(value)
        event_dict[key] = value
    return event_dict

def configure_logging(log_dir: Path = Path("artifacts"), log_name: str = "train.json"):
    """Console = pretty colors, File = JSON."""
    log_dir.mkdir(parents=True, exist_ok=True)
    log_file = log_dir / log_name
    log_level = os.environ.get("LOG_LEVEL", "INFO").upper()

    shared_processors = [
        structlog.stdlib.add_logger_name,
        structlog.stdlib.add_log_level,
        structlog.stdlib.PositionalArgumentsFormatter(),
        structlog.processors.TimeStamper(fmt="iso"),
        structlog.processors.StackInfoRenderer(),
        structlog.processors.format_exc_info,
        structlog.processors.UnicodeDecoder(),
        custom_serializer_processor,
    ]

    structlog.configure(
        processors=shared_processors + [
            structlog.stdlib.ProcessorFormatter.wrap_for_formatter,
        ],
        logger_factory=structlog.stdlib.LoggerFactory(),
        wrapper_class=structlog.stdlib.BoundLogger,
        cache_logger_on_first_use=True,
    )

    console_renderer = structlog.dev.ConsoleRenderer(
        colors=sys.stdout.isatty(),
        exception_formatter=structlog.dev.RichTracebackFormatter(),
    )
    console_handler = logging.StreamHandler(sys.stdout)
    console_handler.setFormatter(
        structlog.stdlib.ProcessorFormatter(
            processors=[
                structlog.stdlib.ProcessorFormatter.remove_processors_meta,
                console_renderer
            ]
        )
    )

```

Sep 24, 2025 21:57

logging.py

Page 2/2

```

    ],
)

file_renderer = structlog.processors.JSONRenderer()
file_handler = logging.FileHandler(log_file, mode="w")
file_handler.setFormatter(
    structlog.stdlib.ProcessorFormatter(
        processors=[
            structlog.stdlib.ProcessorFormatter.remove_processors_meta,
            file_renderer
        ]
    )
)

root_logger = logging.getLogger()
root_logger.addHandler(console_handler)
root_logger.addHandler(file_handler)
root_logger.setLevel(log_level)

logging.getLogger("matplotlib").setLevel(logging.WARNING)
logging.getLogger("PIL").setLevel(logging.WARNING)

return log_file

```

Sep 24, 2025 21:49

config.py

Page 1/1

```

from pathlib import Path
from typing import Tuple
from pydantic import BaseModel
from pydantic_settings import BaseSettings

class ModelSettings(BaseModel):
    input_depth: int = 1
    layer_depths: list[int] = [16, 32]
    blocks_per_stage: list[int] = [1, 1]
    num_classes: int = 10
    dropout: float = 0.2

class DataSettings(BaseModel):
    batch_size: int = 128
    validation_size: int = 10000
    shuffle_buffer: int = 10000

class TrainingSettings(BaseModel):
    """Settings for model training."""
    num_epochs: int = 3
    learning_rate: float = 0.005
    weight_decay: float = 1e-4          # L2 penalty
    log_interval: int = 100

class PlottingSettings(BaseModel):
    """Settings for logging and saving artifacts."""
    output_dir: Path = Path("artifacts")
    figsize: Tuple[int, int] = (5, 3)
    dpi: int = 200

class AppSettings(BaseSettings):
    """Main application settings."""
    debug: bool = False
    random_seed: int = 31415

    model: ModelSettings = ModelSettings()
    data: DataSettings = DataSettings()
    training: TrainingSettings = TrainingSettings()
    plotting: PlottingSettings = PlottingSettings()

def load_settings() -> AppSettings:
    """Load application settings."""
    return AppSettings()

```



Sep 24, 2025 21:36

log.py

Page 3/4

```
ss\u001b[Om=\u001b[35m0.1782\u001b[Om \u001b[36mstep\u001b[Om=\u001b[35m2300\u001b[Om" }  
{ "event": "\u001b[2m2025-09-25T01:33:58.288806Z\u001b[Om [\u001b[32m\u001b[1minfo \u001b[Om] \u001b[1mtrain_iter \u001b[Om] [\u001b[32m\u001b[1mloss\u001b[Om] [\u001b[32m\u001b[1mval_loss\u001b[Om] [\u001b[32m\u001b[1mtrain_loss\u001b[Om] [\u001b[32m\u001b[1mtrain_val_loss\u001b[Om] [\u001b[32m\u001b[1mtrain_acc\u001b[Om] [\u001b[32m\u001b[1mtrain_val_acc\u001b[Om] [\u001b[32m\u001b[1mtrain_mse\u001b[Om] [\u001b[32m\u001b[1mtrain_val_mse\u001b[Om] [\u001b[32m\u001b[1mtrain_rmse\u001b[Om] [\u001b[32m\u001b[1mtrain_val_rmse\u001b[Om] [\u001b[32m\u001b[1mtrain_f1\u001b[Om] [\u001b[32m\u001b[1mtrain_val_f1\u001b[Om] [\u001b[32m\u001b[1mtrain_auc\u001b[Om] [\u001b[32m\u001b[1mtrain_val_auc\u001b[Om] [\u001b[32m\u001b[1mtrain_precision\u001b[Om] [\u001b[32m\u001b[1mtrain_val_precision\u001b[Om] [\u001b[32m\u001b[1mtrain_recall\u001b[Om] [\u001b[32m\u001b[1mtrain_val_recall\u001b[Om] [\u001b[32m\u001b[1mtrain_hf_score\u001b[Om] [\u001b[32m\u001b[1mtrain_val_hf_score\u001b[Om] [\u001b[32m\u001b[1mtrain_mean\u001b[Om] [\u001b[32m\u001b[1mtrain_val_mean\u001b[Om] [\u001b[32m\u001b[1mtrain_std\u001b[Om] [\u001b[32m\u001b[1mtrain_val_std\u001b[Om] [\u001b[32m\u001b[1mtrain_min\u001b[Om] [\u001b[32m\u001b[1mtrain_val_min\u001b[Om] [\u001b[32m\u001b[1mtrain_max\u001b[Om] [\u001b[32m\u001b[1mtrain_val_max\u001b[Om] [\u001b[32m\u001b[1mtrain_avg\u001b[Om] [\u001b[32m\u001b[1mtrain_val_avg\u001b[Om] [\u001b[32m\u001b[1mtrain_sum\u001b[Om] [\u001b[32m\u001b[1mtrain_val_sum\u001b[Om] [\u001b[32m\u001b[1mtrain_diff\u001b[Om] [\u001b[32m\u001b[1mtrain_val_diff\u001b[Om] [\u001b[32m\u001b[1mtrain_ratio\u001b[Om] [\u001b[32m\u001b[1mtrain_val_ratio\u001b[Om] [\u001b[32m\u001b[1mtrain_percent\u001b[Om] [\u001b[32m\u001b[1mtrain_val_percent\u001b[Om] [\u001b[32m\u001b[1mtrain_multiply\u001b[Om] [\u001b[32m\u001b[1mtrain_val_multiply\u001b[Om] [\u001b[32m\u001b[1mtrain_divide\u001b[Om] [\u001b[32m\u001b[1mtrain_val_divide\u001b[Om] [\u001b[32m\u001b[1mtrain_power\u001b[Om] [\u001b[32m\u001b[1mtrain_val_power\u001b[Om] [\u001b[32m\u001b[1mtrain_sqrt\u001b[Om] [\u001b[32m\u001b[1mtrain_val_sqrt\u001b[Om] [\u001b[32m\u001b[1mtrain_log\u001b[Om] [\u001b[32m\u001b[1mtrain_val_log\u001b[Om] [\u001b[32m\u001b[1mtrain_exp\u001b[Om] [\u001b[32m\u001b[1mtrain_val_exp\u001b[Om] [\u001b[32m\u001b[1mtrain_sin\u001b[Om] [\u001b[32m\u001b[1mtrain_val_sin\u001b[Om] [\u001b[32m\u001b[1mtrain_cos\u001b[Om] [\u001b[32m\u001b[1mtrain_val_cos\u001b[Om] [\u001b[32m\u001b[1mtrain_tan\u001b[Om] [\u001b[32m\u001b[1mtrain_val_tan\u001b[Om] [\u001b[32m\u001b[1mtrain_cot\u001b[Om] [\u001b[32m\u001b[1mtrain_val_cot\u001b[Om] [\u001b[32m\u001b[1mtrain_sec\u001b[Om] [\u001b[32m\u001b[1mtrain_val_sec\u001b[Om] [\u001b[32m\u001b[1mtrain_csc\u001b[Om] [\u001b[32m\u001b[1mtrain_val_csc\u001b[Om] [\u001b[32m\u001b[1mtrain_sinh\u001b[Om] [\u001b[32m\u001b[1mtrain_val_sinh\u001b[Om] [\u001b[32m\u001b[1mtrain_cosh\u001b[Om] [\u001b[32m\u001b[1mtrain_val_cosh\u001b[Om] [\u001b[32m\u001b[1mtrain_tanh\u001b[Om] [\u001b[32m\u001b[1mtrain_val_tanh\u001b[Om] [\u001b[32m\u001b[1mtrain_arcsin\u001b[Om] [\u001b[32m\u001b[1mtrain_val_arcsin\u001b[Om] [\u001b[32m\u001b[1mtrain_arcos\u001b[Om] [\u001b[32m\u001b[1mtrain_val_arcos\u001b[Om] [\u001b[32m\u001b[1mtrain_arctan\u001b[Om] [\u001b[32m\u001b[1mtrain_val_arctan\u001b[Om] [\u001b[32m\u001b[1mtrain_arccot\u001b[Om] [\u001b[32m\u001b[1mtrain_val_arccot\u001b[Om] [\u001b[32m\u001b[1mtrain_arcsec\u001b[Om] [\u001b[32m\u001b[1mtrain_val_arcsec\u001b[Om] [\u001b[32m\u001b[1mtrain_arccsc\u001b[Om] [\u001b[32m\u001b[1mtrain_val_arccsc\u001b[Om] [\u001b[32m\u001b[1mtrain_arcsinh\u001b[Om] [\u001b[32m\u001b[1mtrain_val_arcsinh\u001b[Om] [\u001b[32m\u001b[1mtrain_arccosh\u001b[Om] [\u001b[32m\u001b[1mtrain_val_arccosh\u001b[Om] [\u001b[32m\u001b[1mtrain_artanh\u001b[Om] [\u001b[32m\u001b[1mtrain_val_artanh\u001b[Om] [\u001b[32m\u001b[1mtrain_atanh\u001b[Om] [\u001b[32m\u001b[1mtrain_val_atanh\u001b[Om] [\u001b[32m\u001b[1mtrain_acoth\u001b[Om] [\u001b[32m\u001b[1mtrain_val_acoth\u001b[Om] [\u001b[32m\u001b[1mtrain_aacoth\u001b[Om] [\u001b[32m\u001b[1mtrain_val_aacoth\u001b[Om] [\u001b[32m\u001b[1mtrain_alog\u001b[Om] [\u001b[32m\u001b[1mtrain_val_alog\u001b[Om] [\u001b[32m\u001b[1mtrain_aln\u001b[Om] [\u001b[32m\u001b[1mtrain_val_aln\u001b[Om] [\u001b[32m\u001b[1mtrain_erf\u001b[Om] [\u001b[32m\u001b[1mtrain_val_erf\u001b[Om] [\u001b[32m\u001b[1mtrain_erfc\u001b[Om] [\u001b[32m\u001b[1mtrain_val_erfc\u001b[Om] [\u001b[32m\u001b[1mtrain_gammasq\u001b[Om] [\u001b[32m\u001b[1mtrain_val_gammasq\u001b[Om] [\u001b[32m\u001b[1mtrain_gammainc\u001b[Om] [\u001b[32m\u001b[1mtrain_val_gammainc\u001b[Om] [\u001b[32m\u001b[1mtrain_gammaincd\u001b[Om] [\u001b[32m\u001b[1mtrain_val_gammaincd\u001b[Om] [\u001b[32m\u001b[1mtrain_betainc\u001b[Om] [\u001b[32m\u001b[1mtrain_val_betainc\u001b[Om] [\u001b[32m\u001b[1mtrain_betaincd\u001b[Om] [\u001b[32m\u001b[1mtrain_val_betaincd\u001b[Om] [\u001b[32m\u001b[1mtrain_ellipe\u001b[Om] [\u001b[32m\u001b[1mtrain_val_ellipe\u001b[Om] [\u001b[32m\u001b[1mtrain_ellipem\u001b[Om] [\u001b[32m\u001b[1mtrain_val_ellipem\u001b[Om] [\u001b[32m\u001b[1mtrain_ellipk\u001b[Om] [\u001b[32m\u001b[1mtrain_val_ellipk\u001b[Om] [\u001b[32m\u001b[1mtrain_ellipkm\u001b[Om] [\u001b[32m\u001b[1mtrain_val_ellipkm\u001b[Om] [\u001b[32m\u001b[1mtrain_ellipkn\u001b[Om] [\u001b[32m\u001b[1mtrain_val_ellipkn\u001b[Om] [\u001b[32m\u001b[1mtrain_ellipjn\u001b[Om] [\u001b[32m\u001b[1mtrain_val_ellipjn\u001b[Om] [\u001b[32m\u001b[1mtrain_ellipjnm\u001b[Om] [\u001b[32m\u001b[1mtrain_val_ellipjnm\u001b[Om] [\u001b[32m\u001b[1mtrain_ellipjnn\u001b[Om] [\u001b[32m\u001b[1mtrain_val_ellipjnn\u001b[Om] [\u001b[32m\u001b[1mtrain_ellipdn\u001b[Om] [\u001b[32m\u001b[1mtrain_val_ellipdn\u001b[Om] [\u001b[32m\u001b[1mtrain_ellipddm\u001b[Om] [\u001b[32m\u001b[1mtrain_val_ellipddm\u001b[Om] [\u0
```

[illegible]

Sep 22, 2025 23:54

plotting.py

Page 1/1

```

import matplotlib
import matplotlib.pyplot as plt
import structlog

log = structlog.get_logger()

font = {
    # "family": "Adobe Caslon Pro",
    "size": 10,
}

matplotlib.style.use("classic")
matplotlib.rc("font", **font)

def plot_losses(train_losses, val_losses, settings):
    plt.figure(figsize=(8, 5))
    plt.plot(train_losses, label="Train Loss", alpha=0.7)
    steps_per_epoch = 50000 // settings.training.batch_size
    val_x = [i * steps_per_epoch for i in range(1, len(val_losses)+1)]
    plt.plot(val_x, val_losses, label="Validation Loss", marker="o")
    plt.xlabel("Step")
    plt.ylabel("Loss")
    plt.title("Training & Validation Loss")
    plt.legend()
    plt.grid(True)

    settings.plotting.output_dir.mkdir(parents=True, exist_ok=True)
    output_path = settings.plotting.output_dir / "Loss Graph.pdf"
    plt.savefig(output_path)
    log.info("Saved Loss Graph plot", path=str(output_path))

```

Sep 24, 2025 21:33

\_\_init\_\_.py

Page 1/1

```

import optax
import structlog
from flax import nnx
import jax
import pathlib as Path

from .config import load_settings
from .logging import configure_logging
from .model import Classifier
from .training import train, test_evaluation
from .data import load_mnist

def main() -> None:
    settings = load_settings()
    log_file = configure_logging()
    log = structlog.get_logger()
    log.info("Settings loaded", settings=settings.model_dump())
    print(f"Logs are being saved to {log_file}")

    ds_train, ds_val, ds_test = load_mnist(
        batch_size=settings.data.batch_size,
        validation_size=settings.data.validation_size,
    )
    rngs = nnx.Rngs(params=settings.random_seed, dropout=settings.random_seed +
1)
    model = Classifier(
        rngs=rngs,
        input_depth=settings.model.input_depth,
        layer_depths=settings.model.layer_depths,
        blocks_per_stage=settings.model.blocks_per_stage,
        layer_kernel_sizes=[(3,3)] * len(settings.model.layer_depths),
        num_classes=settings.model.num_classes,
        dropout=settings.model.dropout,
    )

    params = nnx.state(model, nnx.Param)
    num_params = sum(p.size for p in jax.tree_util.tree_leaves(params))
    log.info("Model created", num_params=num_params)

    # LR schedule and optimizer
    train_size = 50000
    steps_per_epoch = train_size // settings.data.batch_size
    decay_steps = settings.training.num_epochs * steps_per_epoch
    lr_scheduler = optax.schedules.cosine_decay_schedule(
        init_value=settings.training.learning_rate,
        decay_steps=decay_steps,
    )
    optimizer = nnx.Optimizer(
        model,
        optax.adamw(
            learning_rate=lr_scheduler,
            weight_decay=settings.training.weight_decay,
        ),
        wrt=nnx.Param,
    )

    _ = train(model, optimizer, ds_train, ds_val, rngs)

    # Test
    log.info("Starting final testing...")
    test_loss, test_acc = test_evaluation(model, ds_test, rngs)
    log.info("Final Test Results", test_loss=test_loss, test_acc=test_acc)

```

Sep 24, 2025 21:34

model.py

Page 1/3

```

import jax
import jax.numpy as jnp
from flax import nnx

class Conv2d(nnx.Module):
    def __init__(
        self,
        in_channels: int,
        out_channels: int,
        kernel_size,
        *,
        strides=(1, 1),
        padding="SAME",
        use_bias=True,
        dropout=0.0,
        rngs: nnx.Rngs,
    ):
        self.conv = nnx.Conv(
            in_channels, out_channels, kernel_size,
            strides=strides, padding=padding, use_bias=use_bias, rngs=rngs
        )
        self.dropout = nnx.Dropout(dropout, rngs=rngs) if dropout > 0.0 else None

    def __call__(self, x, *, rngs: nnx.Rngs):
        x = self.conv(x)
        if self.dropout:
            x = self.dropout(x, rngs=rngs)
        return x

class ResidualBlock(nnx.Module):
    """Basic residual block."""
    def __init__(
        self,
        in_channels: int,
        out_channels: int,
        *,
        stride=(1, 1),
        dropout=0.0,
        rngs: nnx.Rngs,
    ):
        self.norm1 = nnx.GroupNorm(num_features=in_channels, num_groups=8, rngs=rngs)
        self.conv1 = Conv2d(
            in_channels, out_channels, (3, 3),
            strides=stride, dropout=dropout, rngs=rngs,
        )
        self.norm2 = nnx.GroupNorm(num_features=out_channels, num_groups=8, rngs=rngs)
        self.conv2 = Conv2d(
            out_channels, out_channels, (3, 3),
            strides=(1, 1), dropout=0.0, rngs=rngs,
        )
        self.shortcut = (
            Conv2d(
                in_channels, out_channels, (1, 1),
                strides=stride, dropout=0.0, rngs=rngs,
            )
            if in_channels != out_channels or stride != (1, 1)
            else None
        )

    def __call__(self, x, *, rngs: nnx.Rngs):
        out = self.norm1(x)
        out = nnx.relu(out)
        out = self.conv1(out, rngs=rngs)

        out = self.norm2(out)

```

Sep 24, 2025 21:34

model.py

Page 2/3

```

        out = nnx.relu(out)
        out = self.conv2(out, rngs=rngs)

        identity = x if self.shortcut is None else self.shortcut(x, rngs=rngs)
        return out + identity

class ResidualStage(nnx.Module):
    def __init__(
        self,
        in_channels: int,
        out_channels: int,
        num_blocks: int,
        dropout: float,
        stride=(1, 1),
        *,
        rngs: nnx.Rngs,
    ):
        self.blocks = []
        self.blocks.append(
            ResidualBlock(
                in_channels,
                out_channels,
                stride=stride,
                dropout=dropout,
                rngs=rngs,
            )
        )
        # Rest keep stride = (1,1)
        for _ in range(1, num_blocks):
            self.blocks.append(
                ResidualBlock(
                    out_channels,
                    out_channels,
                    stride=(1, 1),
                    dropout=dropout,
                    rngs=rngs,
                )
            )

    def __call__(self, x, *, rngs: nnx.Rngs):
        for block in self.blocks:
            x = block(x, rngs=rngs)
        return x

class Classifier(nnx.Module):
    """Residual network classifier for MNIST."""
    def __init__(
        self,
        input_depth: int,
        layer_depths: list[int],
        blocks_per_stage,
        layer_kernel_sizes: list[tuple[int, int]],
        num_classes: int,
        *,
        rngs: nnx.Rngs,
        dropout: float = 0.1,
    ):
        assert len(layer_depths) == len(layer_kernel_sizes), \
            "layer_depths and layer_kernel_sizes must match in length"
        self.init_conv = Conv2d(input_depth, layer_depths[0],
                                kernel_size=layer_kernel_sizes[0],
                                strides=(1, 1), dropout=dropout, rngs=rngs)

        # Build stages
        self.stages = []
        self.blocks_per_stage = blocks_per_stage
        in_ch = layer_depths[0]
        for i, out_ch in enumerate(layer_depths):

```



Sep 24, 2025 21:34

model.py

Page 3/3

```

        stride = (2, 2) if i > 0 else (1, 1) # downsample at stage boundaries
    num_blocks = (self.blocks_per_stage[i]
                  if self.blocks_per_stage is not None
                  else 1)
    stage = ResidualStage(
        in_channels=in_ch,
        out_channels=out_ch,
        num_blocks=num_blocks,
        stride=stride,
        dropout=dropout,
        rngs=rngs,
    )
    self.stages.append(stage)
    in_ch = out_ch

    # Final classifier head
    self.fc = nn.Linear(in_ch, num_classes, rngs=rngs)

def __call__(self, x, *, rngs: nn.Rngs):
    x = self.init_conv(x, rngs=rngs)
    for stage in self.stages:
        x = stage(x, rngs=rngs)

    x = jnp.mean(x, axis=(1, 2))

    # Final logits
    return self.fc(x)

```

Sep 24, 2025 21:57	training.py	Page 1/3
<pre> import jax import jax.numpy as jnp import optax from flax import nnx import matplotlib.pyplot as plt import structlog  from .config import load_settings  log = structlog.get_logger()  def cross_entropy_loss(logits, labels):     return optax.softmax_cross_entropy_with_integer_labels(         logits, labels     ).mean()  def compute_accuracy(logits, labels):     preds = jnp.argmax(logits, axis=-1)     return jnp.mean(preds == labels)  def _loss_fn(model, batch, rngs):     images, labels = batch     labels = labels.ravel()     logits = model(images, rngs=rngs)     loss = cross_entropy_loss(logits, labels)     return loss, logits  @nnx.jit def train_step(model, optimizer, batch, rngs):     grad_fn = nnx.value_and_grad(_loss_fn, has_aux=True)     (loss, logits), grads = grad_fn(model, batch, rngs)     optimizer.update(model, grads)     return loss, logits  @nnx.jit def eval_step(model, batch, rngs):     loss, logits = _loss_fn(model, batch, rngs)     return loss, logits  def evaluate(model, dataset, rngs):     total_loss, total_acc, n_batches = 0.0, 0.0, 0     for batch in dataset.as_numpy_iterator():         loss, logits = eval_step(model, batch, rngs)         _, labels = batch         labels = labels.ravel()         acc = compute_accuracy(logits, labels)         loss = float(loss)         acc = float(acc)         total_loss += loss         total_acc += float(acc)         n_batches += 1     return total_loss / max(n_batches, 1), total_acc / max(n_batches, 1)  def train(model, optimizer, train_ds, val_ds, rngs):     settings = load_settings()     output_dir = settings.plotting.output_dir     output_dir.mkdir(parents=True, exist_ok=True)      log_interval = settings.training.log_interval     num_epochs = settings.training.num_epochs      metrics_history = {         "train_steps": [],         "train_loss": [],         "train_accuracy": [], </pre>		

Sep 24, 2025 21:57	training.py	Page 2/3
<pre>         "val_steps": [],         "val_loss": [],         "val_accuracy": [],     }      global_step = 0      for epoch in range(1, num_epochs + 1):         for batch in train_ds.as_numpy_iterator():             loss, logits = train_step(model, optimizer, batch, rngs)              if global_step % log_interval == 0:                 images, labels = batch                 labels = labels.ravel()                 acc = compute_accuracy(logits, labels)                  loss_val = float(loss)                 acc_val = float(acc)                  metrics_history["train_steps"].append(global_step)                 metrics_history["train_loss"].append(loss_val)                 metrics_history["train_accuracy"].append(acc_val)                  log.info(                     "train_iter",                     epoch=epoch,                     step=global_step,                     loss=loss_val,                     acc=acc_val,                 )                 global_step += 1          val_loss, val_acc = evaluate(model, val_ds, rngs)         metrics_history["val_steps"].append(global_step)         metrics_history["val_loss"].append(val_loss)         metrics_history["val_accuracy"].append(val_acc)          log.info(             "epoch_summary",             epoch=epoch,             step=global_step,             val_loss=val_loss,             val_acc=val_acc,         )      log.info("Training finished. Generating final plots...")     fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))      ax1.set_title("Loss vs. Steps")     ax2.set_title("Accuracy vs. Steps")      ax1.plot(metrics_history["train_steps"], metrics_history["train_loss"], label="train",             alpha=0.7)     ax2.plot(metrics_history["train_steps"], metrics_history["train_accuracy"], label="train",             alpha=0.7)      ax1.plot(metrics_history["val_steps"], metrics_history["val_loss"], label="val",             marker='o', linestyle='--')     ax2.plot(metrics_history["val_steps"], metrics_history["val_accuracy"], label="val",             marker='o', linestyle='--')      ax1.set_xlabel("Step")     ax1.set_ylabel("Loss")     ax1.legend()     ax1.grid(True, which='both', linestyle='--', linewidth=0.5)      ax2.set_xlabel("Step") </pre>		

Sep 24, 2025 21:57

training.py

Page 3/3

```

ax2.set_ylabel("Accuracy")
ax2.legend()
ax2.grid(True, which='both', linestyle='--', linewidth=0.5)

fig.tight_layout()
fig.savefig(output_dir / "final_training_metrics.png", dpi=settings.plotting.dpi)
plt.close(fig)

return metrics_history

def test_evaluation(model, test_ds, rngs):
    settings = load_settings()
    output_dir = settings.plotting.output_dir
    output_dir.mkdir(parents=True, exist_ok=True)

    test_loss, test_acc = evaluate(model, test_ds, rngs)
    log.info("test_results", test_loss=test_loss, test_acc=test_acc)

    test_batch = next(test_ds.as_numpy_iterator())
    _, logits = eval_step(model, test_batch, rngs)
    images, labels = test_batch
    preds = jnp.argmax(logits, axis=-1)

    fig, axs = plt.subplots(5, 5, figsize=(10, 10))
    for i, ax in enumerate(axs.flatten()):
        ax.imshow(images[i, ..., 0], cmap="gray")
        ax.set_title(f"pred={int(preds[i])}, label={int(labels[i])}")
        ax.axis("off")
    fig.tight_layout()
    fig.savefig(output_dir / "test_predictions.png", dpi=settings.plotting.dpi)
    plt.close(fig)

    return test_loss, test_acc

```

Sep 24, 2025 21:57

data.py

Page 1/1

```
import tensorflow_datasets as tfds
import tensorflow as tf

def load_mnist(batch_size=128, validation_size=10000, seed=0):
    ds_train = tfds.load("mnist", split="train", shuffle_files=True, as_supervised=
True)
    ds_test  = tfds.load("mnist", split="test", as_supervised=True)

    def _preprocess(image, label):
        image = tf.cast(image, tf.float32) / 255.0
        label = tf.cast(label, tf.int32)
        return image, label


    ds_train = ds_train.map(_preprocess, num_parallel_calls=tf.data.AUTOTUNE)
    ds_test  = ds_test.map(_preprocess, num_parallel_calls=tf.data.AUTOTUNE)

    ds_val = ds_train.take(validation_size)
    ds_train = ds_train.skip(validation_size)


    ds_train = ds_train.shuffle(10 * batch_size, seed=seed).batch(batch_size).pr
efetch(tf.data.AUTOTUNE)
    ds_val    = ds_val.batch(batch_size).prefetch(tf.data.AUTOTUNE)
    ds_test   = ds_test.batch(batch_size).prefetch(tf.data.AUTOTUNE)

    return ds_train, ds_val, ds_test
```


Loss




Accuracy




pred=2, label=2




pred=0, label=0




pred=4, label=4




pred=8, label=8




pred=7, label=7




pred=6, label=6




pred=0, label=0




pred=6, label=6




pred=3, label=3




pred=1, label=1




pred=8, label=8



pred=0, label=0




pred=7, label=7




pred=9, label=9




pred=8, label=8




pred=4, label=4



pred=5, label=5




pred=3, label=3




pred=4, label=4




pred=0, label=0




pred=6, label=6




pred=6, label=6



pred=3, label=3



pred=0, label=0



pred=2, label=2

